

JOB DESCRIPTION

Junior Data Scientist – FinSight Analytics, Hyderabad

Responsibilities:

- Build and evaluate ML models (classification, regression, clustering)
- Write clean Python code for data wrangling, EDA, and feature engineering
- Collaborate with analysts using SQL to pull and prepare data
- Present findings via dashboards or reports to non-technical stakeholders
- Use version control (Git) for project collaboration

Requirements:

- Proficiency in Python (pandas, NumPy, scikit-learn)
- Working knowledge of SQL
- Familiarity with ML concepts and model evaluation (precision, recall, AUC-ROC)
- Strong analytical and communication skills
- Degree in CS, Statistics, Data Science, or related field
- Nice to have: NLP, time-series analysis, or financial data experience

RESUME

SUMMARY

Final-year Statistics & Data Science student with hands-on experience building ML pipelines on financial and NLP datasets. Comfortable with the full workflow from SQL data extraction to model deployment and stakeholder reporting.

EDUCATION

B.Tech – Statistics & Data Science

IIT Hyderabad | CGPA 8.7 | Expected May 2024

TECHNICAL SKILLS

Python: pandas, NumPy, scikit-learn, XGBoost, NLTK

SQL: PostgreSQL, window functions, CTEs

ML: Classification, Regression, Clustering, NLP (TF-IDF, BERT fine-tuning)

Tools: Git/GitHub, Jupyter, Tableau, FastAPI

EXPERIENCE

Data Science Intern – Zerodha (Bangalore, Jun–Aug 2023, 3 months)

- Built a churn prediction model (XGBoost) on 500K+ user records; improved retention targeting by 18% (AUC-ROC: 0.91)
- Wrote SQL pipelines to pull and clean transactional data from PostgreSQL
- Presented findings weekly to a non-technical product team via Tableau dashboards

PROJECTS

1. Loan Default Prediction (FinTech)

- End-to-end pipeline: feature engineering, SMOTE for class imbalance, tuned Random Forest (precision 0.87, recall 0.83)
- Deployed as REST API using FastAPI; documented on GitHub

2. Sentiment Analysis on Earnings Call Transcripts (NLP)

- Fine-tuned DistilBERT on financial text; correlated sentiment scores with next-day stock movement ($r = 0.62$)

3. Time-Series Forecasting – Nifty 50 Returns

- ARIMA + LSTM hybrid model; RMSE reduced 14% vs baseline

ACHIEVEMENTS

- Winner – IIT-H DataThon 2023 (50+ teams)
- Kaggle: Expert tier, top 8% in two competitions

PUBLICATIONS / TALKS

- Presented at college symposium: "Bias in Financial ML Models"